
Herd Immunity and Learning Externalities in Markets of Adaptive Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A learning agent can satisfy every single-agent stability criterion and still help
2 destabilize the market it operates in. When adaptive agents share an environment
3 that reacts to their decisions, each agent’s retraining reshapes the data every other
4 agent learns from next, a *learning externality* that no agent internalizes. We
5 show its strength is governed not by the number of firms but by the *effective*
6 *number of independent learners*, the leading eigenvalue of the correlation matrix of
7 firms’ response Jacobians, so foundation-model concentration enters the stability
8 condition as a term rather than as a concern: fifty firms fine-tuning one vendor’s
9 model are dynamically close to one large learner. Three closed-form levers follow.
10 A crowding-cadence frontier bounds how often a market can retrain, and closes
11 past a critical crowding level. A herd-immunity threshold makes the market stable
12 if and only if the corrected fraction exceeds $(1 - 1/m_N)/e$ at correction efficacy e ,
13 the imperfect-vaccine coverage law with the systemic modulus as the reproduction
14 number. A Pigouvian wedge prices the externality, and every market of two or more
15 firms over-adapts. The last two levers are substitutes, along a frontier we compute.
16 In a shared-pool order-flow market the mechanism reproduces a published $1.74\times$
17 and $3.16\times$ amplification at two and three firms, bit for bit.

18 1 Introduction

19 A learning agent can pass every stability test we know how to run and still help destabilize the market
20 it operates in.

21 The test we know how to run is a single-agent test. A model is evaluated, stress-tested and certified
22 alone, against a distribution treated as given. When the deployed model reshapes that distribution, the
23 performative prediction literature supplies the correction: repeated retraining converges if and only if
24 a modulus $m = \varepsilon\beta/\gamma$ stays below one, where ε measures how strongly the environment responds
25 and γ is the learner’s own objective curvature (Perdomo et al., 2020). Certify $m < 1$ and the loop is
26 safe.

27 It is not, once the agent has company. Consider dealers whose quotes reshape a common pool
28 of informed flow, recommenders competing for one attention pool, or lenders pricing against one
29 population of borrowers. Each agent’s retraining changes the data every *other* agent will learn from
30 next, and no agent internalizes that effect. We call this a **learning externality**, and it has the property
31 that makes externalities worth studying: every participant behaves rationally, every participant passes
32 its own test, and the interaction amplifies a loop none of them chose. The intuition is not new.
33 Banks that each diversify optimally hold similar portfolios and fail together (Beale et al., 2011;
34 Wagner, 2010), and funds running similar strategies unwound simultaneously in August 2007 with
35 the crowding invisible from any single risk report (Khandani and Lo, 2007). What is new is that the
36 agents now *learn*, which changes both what determines the danger and what can be done about it.

37 **What determines it is not how many firms there are.** Let firm i have a response Jacobian
38 E_i describing how its own deployment reshapes the flow it faces, and let R be the correlation
39 matrix of the vectorized E_i . We show the market’s stability is governed by $m_N = N_{\text{eff}} m_1$ with
40 $N_{\text{eff}} = 1 + \kappa(\lambda_{\max}(R) - 1)$, where κ is the spillover between firms and $\lambda_{\max}(R)$ is what we call the
41 **effective number of independent learners**. A hundred firms perturbing the environment in a hundred
42 orthogonal directions are dynamically one firm. Fifty dealers fine-tuning one vendor’s foundation
43 model are dynamically close to one very large learner, and the market they share inherits that learner’s
44 instability. Decomposing each firm’s response into a shared component and an idiosyncratic one,
45 with s the fraction attributable to a common model, vendor or pretraining corpus, gives $m_N =$
46 $m_1(1 + \kappa s(N - 1))$. **The effective number of learners is the number of independent models,**
47 **not the number of firms.** The quantity a competition regulator measures is the wrong one: fifty
48 equal-share firms have a minimal Herfindahl index and look perfectly competitive, and if all fifty
49 share one vendor they are dynamically a monoculture (Hirschman, 1964).

50 **What can be done about it takes three forms, and two are substitutes.** Because every result is
51 stated in m_N , each inherits the supply chain for free. *Slow down:* firms taking K gradient steps per
52 deployment rather than retraining to convergence face a window $K < K_{\max}(m_N)$ that narrows as
53 crowding rises and closes past a critical crowding level beyond which no retraining frequency is safe.
54 *Correct:* feedback-aware updates (Izzo et al., 2021) stabilize a single learner beyond its boundary,
55 and run inside the market they produce a coverage threshold that is exactly the epidemiological
56 herd-immunity law with the systemic modulus as the reproduction number. *Diversify:* that threshold
57 rises in s , so a market reaches stability along either axis, and the iso-stability frontier between the
58 two is this paper’s most directly usable object. Pricing the externality completes the standard policy
59 triple, and the decentralized equilibrium over-adapts for every $N \geq 2$.

60 **Contributions.** (1) The effective number of independent learners as the reproduction number of a
61 market of adaptive models, turning foundation-model concentration into a term in a stability condition,
62 with two counterexamples invalidating mean-similarity diversity indices. (2) The crowding-cadence
63 frontier and its critical crowding level, composing lazy deployment with multi-agent amplification.
64 (3) A mixed-market herd-immunity theorem, exact at finite correction strength, whose clean form is
65 the imperfect-vaccine coverage law and which carries a critical efficacy. (4) The substitution frontier
66 between model diversity and corrected learning, which exists only once (1) and (3) are expressed
67 in one quantity. (5) The Pigouvian wedge and the over-adaptation corollary, unifying the levers as
68 quantity, technology and price instruments.

69 2 Related work

70 **Monoculture, and why it is static.** Kleinberg and Raghavan (2021) show that when many firms
71 adopt the same algorithm welfare can fall even when the shared algorithm is the better one, because
72 correlated selection destroys the option value of independent evaluation; Bommasani et al. (2022)
73 formalize systemic exclusion when many deployments inherit one model. Both are cross-sectional
74 and measured at one point in time on one cohort. Neither has a retraining loop, so neither can express
75 a stability boundary or a rate, and a market can be statically fine and dynamically fragile at the same
76 shared-model fraction.

77 **Multiplayer performative prediction.** Narang et al. (2023) is the nearest dynamics: multiple
78 learners in a shared decision-dependent environment, with equilibrium concepts and convergence
79 conditions. Those conditions are stated as scalar sensitivity bounds that are blind to the *direction*
80 in which each learner perturbs the environment, which is exactly the object that carries this paper’s
81 content, and the setting has no model-provenance structure and no policy analysis. The single-learner
82 lazy-deployment slope of Mendler-Dünner et al. (2020) is what Theorem 2 composes with multi-agent
83 amplification, a composition not previously made; the corrected update of Izzo et al. (2021) is what
84 Theorem 3 puts into the game, having only ever been run single-learner.

85 **Systemic risk, epidemiology, and the statistic.** Beale et al. (2011) asked exactly this question of
86 banks, showing individually prudent portfolio choices can be collectively destabilizing; their agents
87 do not learn, so the homogeneity is a portfolio overlap rather than an alignment of feedback directions,
88 and there is no second instrument to trade diversity against. Diekmann et al. (1990) is what converts

the epidemiological parallel from analogy into identity: R_0 in a heterogeneous population is *defined* as the spectral radius of a next-generation operator, and m_N is the spectral radius of a linearized joint retraining operator, so the two thresholds are the same mathematical statement and the vaccination law transfers with its refinements. On the statistic, the largest eigenvalue of a correlation matrix as an effective count is standard in random-matrix finance (Laloux et al., 1999; Plerou et al., 2002), signal processing (Roy and Vetterli, 2007) and ecology (Hill, 1973); the novelty is the matrix it is computed from, response Jacobians rather than returns, and the condition it enters. Work reading macroeconomic state off public sentiment (PEBSA, 2024) asks the same question as Section 8 in the opposite measurement regime: there the channel is exogenous, and here it is endogenous, so the signal strengthens as the system nears instability.

3 Setup: private against systemic stability

The base result. Consider N symmetric dealers quoting the same instruments and sharing one pool of informed order flow, coupled by a spillover $\kappa \in [0, 1]$: one dealer tightening its quotes raises the toxic flow routed to every dealer, in proportion κ . Each dealer periodically retrains on the flow it has observed and redeploys. Write m_1 for a single firm’s performative retraining modulus, so a firm in isolation is stable exactly when $m_1 < 1$. The joint retraining map has Jacobian $J = -m_1[(1 - \kappa)I + \kappa \mathbf{1}\mathbf{1}^\top]$, with a common mode at $-m_1 N_{\text{eff}}$, $N_{\text{eff}} = 1 + \kappa(N - 1)$, and $N - 1$ differential modes at $-m_1(1 - \kappa)$. The common mode dominates, so the market crosses into instability a factor N_{eff} earlier than any individual dealer’s own loop would. Measured in a shared-pool simulator with real spreads, inventory and liquidity feedback, amplification is $1.74\times$ at $N = 2$ and $3.16\times$ at $N = 3$ against a linear prediction of 2 and 3 (REFLEX, 2026). This is inherited scaffolding and we do not re-derive it; everything below replaces one object in it.

Two questions, and they are not the same question. Is each agent stable alone, $m_1 < 1$? And is the collection stable when its members interact, $m_N < 1$? A shared environment separates them. Dealer A retrains and redeploys, the pool changes, dealer B observes a changed environment and retrains on it, which changes the pool again, which A then retrains on. The loop closes through the environment, and nothing in A’s validation set contains B. **Agent count is not a descriptive characteristic of a market; it enters the stability condition.** Read as economics, this is an externality, and it is *technological rather than pecuniary* (Buchanan and Stubblebine, 1962): the coupling runs through the data-generating process each learner faces, not through the price at which firms trade with one another. The distinction is load-bearing. Firms affecting each other only through prices impose no market failure, because the price movement is a transfer; firms that reshape the distribution their competitors learn from change the production technology of everyone’s forecast, which is a genuine efficiency loss. We call m_N the **feedback reproduction number**: the factor by which a common perturbation is amplified per retraining round.

Objects and standing assumptions. Firm i has a response Jacobian $E_i \in \mathbb{R}^{d \times d}$. The *alignment matrix* $R = (r_{ij})$ collects $r_{ij} = \langle \text{vec } E_i, \text{vec } E_j \rangle / (\|E_i\|_F \|E_j\|_F)$, a correlation matrix of feedback directions rather than of returns; it is positive semidefinite with unit diagonal. The generalization is the single substitution of R for $\mathbf{1}\mathbf{1}^\top$.

Assumption 1 (Standing). (A1) All spectral statements linearize the joint retraining map around the joint equilibrium, with the simulator as the nonlinear check. (A2) Theorem 1 is exact for equal moduli $m_i = m_1$; unequal moduli give a two-sided bound. (A3) Firms retrain on a common clock at a common cadence K . (A4) Firms couple through one pool with a scalar spillover κ . (A5) The common mode binds, which holds when R has nonnegative entries, where Perron-Frobenius places the leading eigenvector in the nonnegative orthant. Shared vendors and shared corpora produce positive alignment, so the regime we study satisfies it; Section 5 does not need it at all.

4 Result 1: the effective number of independent learners

Counting firms is wrong, and the correction is where monoculture and supply-chain content enter. $\lambda_{\max}(R)$ runs from 1, all responses orthogonal, to N , monoculture.

Theorem 1 (Alignment sets the reproduction number). *Under Assumption 1 with equal moduli and heterogeneous response directions, the joint Jacobian is $J = -m_1[(1 - \kappa)I + \kappa R]$ and the market’s*

Table 1: Anchors of Theorem 1. The simplex is the sharper counterexample to mean-based diversity: it *minimizes* mean alignment yet exceeds orthogonality’s $\lambda_{\max}$.

Configuration	$\lambda_{\max}(R)$	N_{eff}	Reading
Monoculture, $R = \mathbf{1}\mathbf{1}^\top$	N	$1 + \kappa(N - 1)$	the inherited law
Orthogonal, $R = I$	1	1	a hundred firms are dynamically one firm
Simplex, $r_{ij} = -1/(N-1)$	$N/(N-1)$	$1 + \kappa/(N-1)$	maximal diversity, and the mode labels swap
Supply chain, share s	$1 + s(N - 1)$	$1 + \kappa s(N - 1)$	the number of independent <i>models</i>

140 *spectral radius is*

$$m_N = N_{\text{eff}} m_1, \quad N_{\text{eff}} = 1 + \kappa(\lambda_{\max}(R) - 1), \quad (1)$$

141 *so the market is stable if and only if $m_N < 1$.*

142 Setting $R = \mathbf{1}\mathbf{1}^\top$ recovers the symmetric multi-dealer law of REFLEX (2026) exactly, so the
 143 published result is not generalized away but located: it is the monoculture corner, where every firm’s
 144 feedback points the same way. Three anchors follow, each checked independently (Table 1). $N_{\text{eff}} \geq 1$
 145 always, so interaction never stabilizes a market below what its members achieve alone.

146 **The supply chain.** Decompose $E_i = \sqrt{s} E_{\text{shared}} + \sqrt{1-s} \Xi_i$, where $s \in [0, 1]$ is the fraction
 147 attributable to a shared foundation model, vendor or pretraining corpus. Alignments concentrate at
 148 $r_{ij} \rightarrow s$ with $O(1/d)$ fluctuation and an explicit failure probability, so $N_{\text{eff}} \rightarrow 1 + \kappa s(N - 1)$ and
 149 every result below inherits s for free. Fifty dealers fine-tuning one vendor’s model are dynamically
 150 $\sim 1 + 49\kappa$ learners at $s \approx 1$ and barely more than one at $s \approx 0$.

151 **Why the naive diversity index fails, twice.** Mean alignment understates *clustered* alignment: three
 152 tightly aligned firms among ten otherwise-orthogonal ones give $N_{\text{eff}} = 2.60$ at $\kappa = 0.8$ against 1.48
 153 for the uniform configuration with the identical mean, an understatement of $1.76\times$. Concretely, a
 154 market at $m_1 = 0.5$ runs $m_N = 1.30$ and is unstable while the mean index reports 0.74 and calls it
 155 safe with margin. The simplex makes the sharper point: it minimizes mean alignment at $-1/(N - 1)$,
 156 yet its $\lambda_{\max} = N/(N - 1)$ exceeds orthogonality’s 1, so the mean *does not order configurations*
 157 *correctly* and fails as a ranking rather than only as a level. We prove a signed version of the failure:
 158 $N_{\text{eff}} \geq 1 + \kappa(N - 1)\bar{r}$ for every R , so the mean index always under-states risk and never over-states
 159 it. The rule the paper follows is that all stability claims use the spectral form, never a mean.

160 **Naming.** We do not present the *quantity* as new, only its role. Spectral summaries of a correlation
 161 matrix have a long history as effective counts (Hill, 1973; Roy and Vetterli, 2007; Laloux et al.,
 162 1999). What is new is the object it is computed from, response Jacobians rather than returns, and the
 163 condition it enters.

164 5 Result 2: the crowding-cadence frontier

165 The first lever is quantity regulation, and it is the one a firm can pull alone: retrain less often. A learner
 166 taking only K gradient steps per deployment realizes outer-map slope $\mu(K) = -m + c^K(1 + m)$,
 167 where $c = 1 - \eta\gamma$ is the inner per-step contraction (Mendler-Dünner et al., 2020). Because c is built
 168 from the learner’s own objective curvature and step size, it does not depend on how many competitors
 169 it has, and that independence is what makes the composition below work. We verify it by measuring
 170 the realized contraction inside the joint market rather than asserting it.

171 **Theorem 2** (Systemic cadence window). *In the N -firm market with synchronous K -step retraining,*
 172 *the joint deployment map is $c^K I + (1 - c^K)J$, so the binding mode has slope $\mu_N(K) = -m_N +$*
 173 *$c^K(1 + m_N)$ and the market is stable if and only if*

$$K < K_{\max} = \frac{\ln((m_N - 1)/(m_N + 1))}{\ln c}, \quad (2)$$

174 *with $K_{\max}$ infinite when $m_N \leq 1$ and strictly decreasing in m_N .*

175 The slope on the mode carrying $\lambda_i(R)$ is strictly decreasing in λ_i , which does two things at once: it
 176 puts every slope at or below $c^K < 1$, so the upper side of the stability constraint never binds at any




Figure 1: Left: the (N, s) phase diagram, measured m_N against the predicted boundary over 48 cells (Theorem 1). Right: the cadence window over 175 cells of the (N_{eff}, K) grid against predicted K_{max} , with s contours and $c = 0.8$ (Theorem 2); the realized window is $\lfloor K_{\text{max}} \rfloor$ while the plotted frontier is continuous. Both panels are dry runs in the reference environment (Section 9).

cadence, and it places the binding mode at λ_{max} whatever the sign pattern of R , so this section needs no analogue of (A5).

Critical crowding. The same inequality reads $m_N < (1 + c^K)/(1 - c^K)$, and since K is a positive integer the most permissive case is $K = 1$. A stabilizing cadence exists at all if and only if $m_N < (1 + c)/(1 - c)$, a factor of 9 at $c = 0.8$. Past that the market is unstable at every retraining frequency. This is not a second result but the frontier at its smallest admissible cadence, which is also why the laziest retrainer is the hardest to destabilize.

What the supply chain does to the window. Substituting Result 1 makes K_{max} a function of vendor concentration at a fixed number of firms. At $m_1 = 0.15$, $\kappa = 0.8$, $N = 30$, $c = 0.8$, the window runs $K_{\text{max}} = 20.68$ at $s = 0.25$, 5.28 at $s = 0.5$ and 2.53 at $s = 1$. Holding competitors fixed and raising vendor concentration cuts every incumbent’s retraining budget by a factor of eight. The externality in its sharpest operational form is therefore not merely that **your competitor’s entry consumes your retraining budget** but that **your competitor’s choice of vendor does too**, without their entering at all and without either firm doing anything wrong.

What the lever costs. Cadence buys stability with model staleness, and the exchange rate is the same quantity on both sides: after a retraining round the deployed model’s gap to its own best response is exactly c^K times what it was, so the c^K that widens the stability margin is the staleness. We quantify it rather than presenting it as free.

6 Result 3: herd immunity, and diversity as its substitute

The second lever is a technology mandate, and unlike cadence it is not something a firm can pull alone to any effect. A *corrected* firm runs the feedback-aware update (Izzo et al., 2021; Miller et al., 2021), so its retraining is governed by the objective curvature γ_{PO} in place of the cobweb and its modulus is scaled by $\theta = \gamma/\gamma_{\text{PO}} < 1$. What is new is running it inside the game: a mixed market of N_b blind firms and $N - N_b$ corrected ones, corrected fraction $\rho = 1 - N_b/N$. Correction changes a firm’s gain and not its response direction, so R is untouched.

Theorem 3 (Mixed-market stability). *With supply-chain alignment and two correction levels, the joint Jacobian’s symmetric congruence is a diagonal matrix plus a positive rank-one term, so its spectrum solves a secular equation. Two distinct moduli collapse it to a quadratic, and for $1 \leq N_b \leq N - 1$ the radius is its larger root,*

$$\rho(J) = \frac{1}{2}(P + \sqrt{P^2 - 4Q}), \quad P = (1 - \tilde{\kappa})m_1(1 + \theta) + \tilde{\kappa}m_1(N_b + \theta N_c), \quad Q = \theta m_1^2(1 - \tilde{\kappa})N_{\text{eff}}, \quad (3)$$

with $\tilde{\kappa} = \kappa s$ and $N_c = N - N_b$. An empty block needs the single-block form instead; the quadratic otherwise leaves a phantom root that can exceed the true radius.

208 N_{eff} appearing inside Q is why every statement below carries s : the mixed market inherits Result 1's
 209 crowding through the product term rather than by stipulation.

210 **The strong-correction limit, and it errs in the unsafe direction.** As $\gamma_{\text{PO}} \rightarrow \infty$ the constant term
 211 vanishes and the unstable cycle runs through blind firms only, giving $N_b < N_c(s) = 1 + (1/m_1 -$
 212 $1)/(\kappa s)$. This is the memorable criterion, and taken alone it is **optimistic rather than conservative**.
 213 The radius is nondecreasing in θ by Perron-Frobenius on an entrywise-nonnegative matrix, so the
 214 limit under-states it: on random draws the limit calls 11.8% of configurations stable that are unstable
 215 at every finite correction strength. The exact root is therefore the criterion and the limit is its corollary,
 216 quoted with its error direction rather than on its own. The same monotonicity carries the reassuring
 217 half: **correction never backfires**, so no configuration exists in which un-blinding a firm destabilizes
 218 the market.

219 **The collapse, and its honest generalization.** At $\kappa = s = 1$ the limit reads $\rho^* = 1 - 1/m_N$,
 220 **exactly the epidemiological herd-immunity threshold** $1 - 1/R_0$ with the systemic modulus as the
 221 reproduction number. A market of 10 firms at $m_N = 2.5$ needs 60% of them un-blinded, *if correction*
 222 *is perfect*. It is not, and at that same corner the quadratic degenerates so that the stability condition
 223 solves exactly for

$$\rho^*(e) = \frac{1 - 1/m_N}{e}, \quad e = 1 - \gamma/\gamma_{\text{PO}}, \quad (4)$$

224 with e the **correction efficacy**. That is the standard coverage requirement for an *imperfect vaccine*,
 225 of which the clean law is the perfect-efficacy corner. The correspondence does not merely survive the
 226 refinement, it predicts it: a structural analogy transfers the refinements of the law it borrows, and a
 227 decorative one does not. The clean $1 - 1/R_0$ form assumes homogeneous mixing (Fine et al., 2011),
 228 and the collapse here assumes the $\kappa = s = 1$ corner for precisely that reason.

229 **Corollary 4** (Critical efficacy). $\rho^*(e)$ exceeds one, so no corrected fraction works at all, exactly
 230 when $\theta > 1/m_N$. Correction is a usable lever only if $\gamma_{\text{PO}}/\gamma > m_N$.

231 At $m_N = 2.5$ the corrected update must deliver a $2.5 \times$ curvature improvement or un-blinding cannot
 232 stabilize the market whatever fraction is treated. This is the exact structural parallel of Section 5's
 233 critical crowding: each lever has a regime past which it stops working, and knowing both is what
 234 makes trading them off an honest exercise rather than an extrapolation.

235 **The synthesis.** ρ^* is increasing in s , strictly wherever positive, so a market reaches stability along
 236 either axis. At $N = 20$, $m_1 = 0.15$, $\kappa = 0.8$: a monoculture ($s = 1$) has $N_c = 8.08$ and needs
 237 $\rho^* = 0.596$, that is 12 corrected firms; at $s = 0.5$, $N_c = 15.17$ and $\rho^* = 0.242$, five firms; at $s = 0.2$,
 238 $N_c = 36.42$ and $\rho^* = 0$, so the market is stable with no corrected agents at all. The last column is
 239 the one a regulator acts on, and it is $N - \lceil N_c \rceil + 1$ rather than $\lceil \rho^* N \rceil$; the two differ by one firm
 240 exactly when N_c is an integer, which at N in the tens is several percentage points of coverage.

241 **The economic reading.** Correction is a public good. An un-blinded firm captures a private benefit
 242 while the stability it contributes accrues to everyone, so the market free-rides below the threshold and
 243 does not reach ρ^* unaided (Bergstrom et al., 1986). Model diversity has the same structure: a firm
 244 choosing an idiosyncratic vendor pays for it privately and supplies systemic stability for free. **The**
 245 **substitution is therefore between two goods that are both under-supplied**, not between a good
 246 and a bad. That is what makes the frontier something a regulator can move along rather than a menu
 247 of two policies plotted on shared axes.

248 7 Result 4: the Pigouvian wedge

249 The third lever is the price instrument. Inside the stable region the binding mode is an AR(1) with
 250 coefficient $-m_N$ driven by noise of variance σ^2 , so its stationary variance is $V(m_N) = \sigma^2/(1 - m_N^2)$.
 251 The choice is defended rather than assumed: stationary variance is the smooth proxy for divergence
 252 risk, finite everywhere strictly inside the stable region and divergent exactly at the boundary the rest
 253 of the paper is about, and nothing below uses more than V being positive, increasing and convex.
 254 The coefficient carries a sign the variance does not see, and it is not lost information: the lag-1
 255 autocorrelation is $-m_N$ and therefore *negative*, so a crowded market's common mode oscillates
 256 rather than persists, which is the observable Section 8 reads.




Figure 2: Left: stability against corrected fraction at $N = 20$, $m_1 = 0.15$, $\kappa = 0.8$, $s = 1$, across efficacies; the gap between the exact threshold and the perfect-correction limit is a result, not an error bar. Right: **the substitution frontier**, the (ρ, s) iso-stability curve, measured against predicted, matching at 18 of 18 points. Both are dry runs (Section 9).

Each firm picks an adaptation aggressiveness a_i , meaning cadence, learning rate or responsiveness, buying concave tracking benefit $B(a_i)$ through $m_i = \mu(a_i)$. Firm i bears weight w_i on common-mode variance and the planner bears W , with $W > \sum_i w_i$ because every trade has a client on the other side who bears execution-quality variance while no dealer's objective contains it. That is an assumption, stated as one, and the results turn on its *sign* and never its size.

Lemma 5 (Marginal crowding share). *At the symmetric point, $\partial m_N / \partial m_i = N_{\text{eff}} v_i^2$ where v is the leading eigenvector of the coupling matrix. The shares sum to N_{eff} rather than to one, and on exchangeable R each equals N_{eff}/N .*

Both readings matter. The *sum* is N_{eff} , so the market's total sensitivity to a uniform increase in aggressiveness is amplified by exactly the effective learner count of Result 1. The *individual* share is N_{eff}/N , the arithmetic shape of every commons problem. The naive guess here is $1/N$ by analogy with the commons, and it is wrong by the factor N_{eff} , which at $N = 20$, $\kappa = 0.8$, $s = 1$ is 16.2.

Theorem 6 (The wedge). *At a symmetric profile the private and social first-order conditions are the same equation with a different price on crowding, and the marginal cost a firm ignores is*

$$t^* = (W - w_i) V'(m_N) \frac{N_{\text{eff}}}{N} \mu'(a), \quad V'(m_N) = \frac{2\sigma^2 m_N}{(1 - m_N^2)^2}. \quad (5)$$

A per-unit fee t^* on aggressiveness makes the private condition coincide with the social one, so the taxed decentralized equilibrium implements the symmetric social optimum.

With symmetric weights the private condition ignores the fraction $(N - 1)/N$ of the marginal variance cost borne by other firms, plus the entire client-side exposure. t^* is strictly increasing in N , κ , s and m_N , and diverges at the boundary like $(1 - m_N)^{-2}$. The rate is the economically loaded part: a regulator who sets a fee from a comfortable-looking market and holds it fixed is setting it far too low by the time the market is close to instability.

Corollary 7 (Over-adaptation). *The decentralized equilibrium over-adapts relative to the social optimum for every $N \geq 2$, strictly so even with client exposure switched off, and the distortion grows as the market crowds. At $N = 1$ with no client exposure the wedge is exactly zero.*

The provenance channel. m_N depends on s , so the wedge has a second channel: $\partial m_N / \partial a_i = (N_{\text{eff}}/N) \mu'(a_i)$ against $\partial m_N / \partial s = m_1 \kappa (N - 1)$. The same wedge therefore prices **shared-model adoption** directly, turning the monoculture externality into a fee rather than a warning. The contrast explains the adoption dynamic in one line. The aggressiveness channel's N_{eff}/N factor is bounded and tends to κ from above, so a firm's marginal footprint through its own cadence does not grow without bound. The provenance channel is **linear in N and does not decay**, while a firm's private quality gain from adopting the market-leading vendor does not grow with market size at all. Unregulated adoption therefore runs toward the unstable configuration rather than away from it, which is why the

Table 2: Panel coverage and status. Panel 1 is external validation against a prior published run, not a self-consistency check.

Panel	Tests	Status	Outcome
1. Amplification	Section 3	measured	$1.7428\times, 3.1567\times$; relative error 0.00
2. (N, s) phase diagram	Thm 1	dry run	48 cells, max error 7.1×10^{-15}
2b. Clustered companion	Thm 1	dry run	measured m_N 1.30; mean index says 0.74
3. Cadence frontier	Thm 2	dry run	175 cells, zero disagreements
4. Herd immunity	Thm 3	dry run	thresholds 12/14/16/20 firms across efficacy
5. Substitution frontier	Thm 3	dry run	18 of 18 points exact
6. Over-adaptation	Thm 6	dry run	12 configurations, zero contradicting Cor. 7

diversity floor of Section 6 is a different instrument from the cadence cap of Section 5 rather than a restatement of it. Sections 5, 6 and 7 are respectively quantity regulation, a technology mandate or structural remedy, and the price instrument; choosing among them under uncertainty is Weitzman’s question (Weitzman, 1974), and we supply the objects needed to pose it rather than answering it.

8 Supervision from public prices

Every quantity above is hidden from the regulator these results address, and near the boundary the system tells on itself: as m_N approaches one the leading principal component of public quotes has lag-1 autocorrelation approaching one in magnitude, with the sign negative because the common mode oscillates rather than persists, and variance share growing like $1/(1 - m_N)$, so a supervisor can estimate distance to instability from public prices alone with no access to any firm’s model, data or code. The estimator sharpens exactly as the market approaches the boundary it watches for. Its consistency proof, central limit theorem and detection sample complexity are deferred, and the real-data panel in the appendix is consistency evidence with a placebo, never identification.

9 Experiments

Six panels, one per result, and each panel’s evidential status is stated in its own row because the six are not the same kind of evidence. A panel marked **measured** ran in the order-flow simulator: a genuine N -dealer market with informed flow, spread and inventory, reducing bit for bit to the single-dealer market at $N = 1$. A panel marked **dry run** ran in a linearized reference environment realizing the model’s response geometry and nothing else. A dry run establishes that the closed forms govern the realized dynamics, which no amount of algebra checking covers, and it establishes nothing about a market. Stability in every panel is estimated from the trajectory’s asymptotic growth rate under the actual retraining map rather than read off an eigensolve, so *measured* always means measured from dynamics. All runs are CPU-only and deterministic from (config, seed), with common random numbers on paired probes; every closed form carries a numerical certificate that fails loudly, and the harness exits nonzero on any disagreement rather than producing a quiet figure. The simulator is described generically here pending an artifact release; code, derivations and certificates are available to reviewers as anonymized supplementary material and are released publicly on acceptance.

The one measured panel. Panel 1 reproduces a prior published run in that run’s own simulator, measuring common-mode moduli 0.7856, 1.3692 and 2.4799 at $N = 1, 2, 3$, hence amplification $1.7428\times$ and $3.1567\times$ against a published $1.74\times$ and $3.16\times$ at relative error 0.00. **The gap to the linear prediction is content, not error:** the market runs -12.9% at $N = 2$ and $+5.2\%$ at $N = 3$, with *opposite signs*, which is what a nonlinear market with saturating flow does around a linearization. The reference environment reproduces the prediction exactly because it omits precisely the nonlinearity and saturation that produce the two-signed departure. The differential mode measures 3.4×10^{-3} against a theoretical zero at $\kappa = 1$, so the instability is purely common-mode, which is the mechanism Theorem 1 generalizes.

Why the rest are not measured, stated plainly. A heterogeneous-response port of the order-flow simulator was designed and built, and it is exact: at flat exposure profiles it reproduces the unmodified simulator at every N from 1 to 6. Its acceptance gate then failed. The simulator’s liquidity and




Figure 3: Over-adaptation and the wedge. Left: decentralized against socially optimal aggressiveness across N at $\chi = 1$, $s = 1$, with the gap shaded. Right: t^* against N on a log scale. Dry run; the boundary divergence rate fits $(1 - m_N)^{-2}$ with log-log slope -2.000 .

price-impact channels stay shared whatever the response profiles are, and isolating that residual shows it reaching 17.7 percentage points across the separation range at $N = 2$, larger than several of the effects these panels resolve; at $N \geq 4$ the finite-difference probe is no longer measuring a local slope at all. We drew no figure from that port and moved no status, which is what the gate is for. Panel 6 is a dry run for a different reason: the order-flow simulator carries no aggressiveness choice variable and no welfare object, so no measured counterpart of it exists to be built.

Panel 6. Both first-order conditions are solved by bisection with m_N read from realized dynamics rather than from the closed form, so the over-adaptation verdict is a property of the simulated market rather than of the algebra. The gap widens from 29.0% at $N = 2$ to 119.9% at $N = 20$, the fee implements the social optimum to 1.7×10^{-13} , and the degenerate $N = 1$ case sits at exactly zero, as Corollary 7 requires.

10 Limitations and conclusion

Limitations. The analysis is linearized around the joint equilibrium (A1), and panel 1 shows what that costs in a real market. **The heterogeneous sweep is closed-form agreement, not measurement,** and this is the largest gap in the work: panels 2 through 6 establish that the closed forms govern realized dynamics in a reference environment with no informed flow, no spread and no inventory, and establish nothing about a real market. Routing the liquidity and impact channels through the response profiles is the identified repair. Theorem 1 is exact for equal moduli (A2); the heterogeneous-modulus case is a two-sided bound whose tightness is measured rather than asserted, and share-weighted alignment for unequal-sized firms, the fully heterogeneous block-secular reduction, and the sharp operator-norm concentration rate are named extensions rather than results. The $1 - 1/R_0$ collapse additionally assumes the $\kappa = s = 1$ corner, three or more correction levels return a genuine secular equation rather than a quadratic, and correction that moves a firm's response direction rather than only its gain is not modelled. Deployment is synchronous on a common cadence (A3); heterogeneous clocks break the eigenvector sharing Theorem 2 uses. Theorem 6 prices instability through stationary variance, a modeling choice. Neither R nor s is observable to an outsider, and Section 8 is the designed answer with its real-data leg deferred.

Conclusion. **Individual model evaluation cannot certify a market of models.** The quantity separating the two questions is the effective number of independent learners, a spectral property of how firms' response directions align, and no evaluation of a single model can see it. A firm can pass every stability test available to it and still be one of the reasons its market fails, without doing anything wrong and without observing the firms it is coupled to. A regulator already drafting systemic-risk obligations for shared general-purpose models has a threshold on model scale but no mechanism connecting model sharing to a measurable systemic quantity. This paper supplies one, and with it two substitutable instruments and a price: retrain less often, share fewer models or correct more of them, or pay for the crowding. The substitution frontier is the object a regulator would act on, because the instruments are interchangeable at a computable rate rather than merely both available.

References

- N. Beale, D. G. Rand, H. Battey, K. Croxson, R. M. May, and M. A. Nowak. Individual versus systemic risk and the regulator’s dilemma. *PNAS*, 108(31):12647–12652, 2011.
- T. Bergstrom, L. Blume, and H. Varian. On the private provision of public goods. *Journal of Public Economics*, 29(1):25–49, 1986.
- R. Bommasani, K. A. Creel, A. Kumar, D. Jurafsky, and P. Liang. Picking on the same person: Does algorithmic monoculture homogenize outcomes? In *NeurIPS*, 2022.
- J. M. Buchanan and W. C. Stubblebine. Externality. *Economica*, 29(116):371–384, 1962.
- O. Diekmann, J. A. P. Heesterbeek, and J. A. J. Metz. On the definition and the computation of the basic reproduction ratio R_0 in models for infectious diseases in heterogeneous populations. *Journal of Mathematical Biology*, 28(4):365–382, 1990.
- P. Fine, K. Eames, and D. L. Heymann. Herd immunity: A rough guide. *Clinical Infectious Diseases*, 52(7):911–916, 2011.
- M. O. Hill. Diversity and evenness: A unifying notation and its consequences. *Ecology*, 54(2):427–432, 1973.
- A. O. Hirschman. The paternity of an index. *American Economic Review*, 54(5):761–762, 1964.
- Z. Izzo, L. Ying, and J. Zou. How to learn when data reacts to your model: Performative gradient descent. In *ICML*, 2021.
- A. E. Khandani and A. W. Lo. What happened to the quants in August 2007? *Journal of Investment Management*, 5(4):29–78, 2007.
- J. Kleinberg and M. Raghavan. Algorithmic monoculture and social welfare. *PNAS*, 118(22):e2018340118, 2021.
- L. Laloux, P. Cizeau, J.-P. Bouchaud, and M. Potters. Noise dressing of financial correlation matrices. *Physical Review Letters*, 83(7):1467–1470, 1999.
- C. Mendler-Dünnér, J. C. Perdomo, T. Zrnic, and M. Hardt. Stochastic optimization for performative prediction. In *NeurIPS*, 2020.
- J. P. Miller, J. C. Perdomo, and T. Zrnic. Outside the echo chamber: Optimizing the performative risk. In *ICML*, 2021.
- A. Narang, E. Faulkner, D. Drusvyatskiy, M. Fazel, and L. J. Ratliff. Multiplayer performative prediction: Learning in decision-dependent games. *JMLR*, 24, 2023.
- Predicting economic behavior via sentiment analysis. *International Journal of Engineering and Computer Science*, 13(12), 2024. <https://doi.org/10.18535/ijecs/v13i12.4950>.
- J. C. Perdomo, T. Zrnic, C. Mendler-Dünnér, and M. Hardt. Performative prediction. In *ICML*, 2020.
- V. Plerou, P. Gopikrishnan, B. Rosenow, L. A. N. Amaral, T. Guhr, and H. E. Stanley. Random matrix approach to cross correlations in financial data. *Physical Review E*, 65(6):066126, 2002.
- Reflexive equilibrium fixed-point learning for endogenous exchanges. *arXiv:2608.16155*, 2026. <https://doi.org/10.48550/arXiv.2608.16155>.
- O. Roy and M. Vetterli. The effective rank: A measure of effective dimensionality. In *EUSIPCO*, 2007.
- W. Wagner. Diversification at financial institutions and systemic crises. *Journal of Financial Intermediation*, 19(3):373–386, 2010.
- M. L. Weitzman. Prices vs. quantities. *Review of Economic Studies*, 41(4):477–491, 1974.

A Technical appendices

Complete proofs of every numbered result in the body, the numerical certificates and their tolerances, the full experimental specifications, and the supervision panel with its placebo are provided in the supplementary material. Every closed form in the body is checked by an assertion-based certificate that fails loudly; 525 assertions across seven files pass at submission time.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state exactly the four closed-form results proved in Sections 4–7 and the one externally validated measurement in Section 9, and they separate the two evidential statuses. The abstract quotes only the measured replication and never a dry run.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 10 is a dedicated limitations paragraph. It names the linearization (A1) with its measured cost, states that five of six panels are reference-environment dry runs rather than market measurements and why the port that would have changed that failed its gate, and lists the equal-moduli restriction, the homogeneous-mixing corner, the synchronous-cadence assumption and the unobservability of R and s .

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Assumption 1 collects the five standing assumptions (A1)–(A5) and each theorem references it. All theorems, lemmas and corollaries are numbered and cross-referenced. Proof sketches appear inline in Sections 4–7 and complete proofs are in the supplementary material, as are the numerical certificates with their tolerances.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Every panel is CPU-only, numpy-only and deterministic from a (config, seed) pair, and Section 9 states the protocol rules and every worked-example parameter, $c = 0.8$ included. The reference environment is specified by its acceptance test, that it reproduces the homogeneous market at $R = 11^\top$, which is checked before any number is taken from it.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

519 (d) We recognize that reproducibility may be tricky in some cases, in which case
520 authors are welcome to describe the particular way they provide for reproducibility.
521 In the case of closed-source models, it may be that access to the model is limited in
522 some way (e.g., to registered users), but it should be possible for other researchers
523 to have some path to reproducing or verifying the results.

524 5. Open access to data and code

525 Question: Does the paper provide open access to the data and code, with sufficient instruc-
526 tions to faithfully reproduce the main experimental results, as described in supplemental
527 material?

528 Answer: [No]

529 Justification: The public repository names an author, so it is withheld at submission to
530 preserve the blind. Code, derivations and all seven certificate files are provided to reviewers
531 as anonymized supplementary material, and the repository is linked in the camera-ready
532 version; Section 9 states this arrangement. No dataset is used: every panel is simulated and
533 deterministic from a (config, seed) pair.

534 Guidelines:

- 535 • The answer [N/A] means that paper does not include experiments requiring code.
- 536 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/
537 public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 538 • While we encourage the release of code and data, we understand that this might not
539 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
540 including code, unless this is central to the contribution (e.g., for a new open-source
541 benchmark).
- 542 • The instructions should contain the exact command and environment needed to run to
543 reproduce the results. See the NeurIPS code and data submission guidelines ([https://
544 neurips.cc/public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 545 • The authors should provide instructions on data access and preparation, including how
546 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 547 • The authors should provide scripts to reproduce all experimental results for the new
548 proposed method and baselines. If only a subset of experiments are reproducible, they
549 should state which ones are omitted from the script and why.
- 550 • At submission time, to preserve anonymity, the authors should release anonymized
551 versions (if applicable).
- 552 • Providing as much information as possible in supplemental material (appended to the
553 paper) is recommended, but including URLs to data and code is permitted.

554 6. Experimental setting/details

555 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
556 rameters, how they were chosen, type of optimizer) necessary to understand the results?

557 Answer: [Yes]

558 Justification: Section 9 and the figure captions state the grid, the parameters (m_1 , κ , s ,
559 N , c , efficacy) and the estimator: stability is read from the trajectory's asymptotic growth
560 rate under the actual retraining map by power iteration, not from an eigensolve. Full
561 specifications are in the supplementary material.

562 Guidelines:

- 563 • The answer [N/A] means that the paper does not include experiments.
- 564 • The experimental setting should be presented in the core of the paper to a level of detail
565 that is necessary to appreciate the results and make sense of them.
- 566 • The full details can be provided either with the code, in appendix, or as supplemental
567 material.

568 7. Experiment statistical significance

569 Question: Does the paper report error bars suitably and correctly defined or other appropriate
570 information about the statistical significance of the experiments?

Answer: [No]

Justification: The panels are deterministic rather than stochastic, so error bars are not the right summary. Agreement is reported instead as the worst-case departure from the closed form over the whole grid (for example max absolute error 7.1×10^{-15} over 48 cells, zero disagreements over 175 cells), which is a stronger statement than a confidence interval over seeds would be. Every certificate carries a stated tolerance and fails loudly.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiments are CPU-only and numpy-only. The full panel suite and all seven certificate files run in under two minutes on a single commodity laptop core, with no GPU, no accelerator and no distributed compute. No unreported large-scale preliminary experiments were run.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work conforms to the NeurIPS Code of Ethics. It is a theoretical and simulation study with no human subjects, no personal data and no released model.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 10 states the policy reading directly: the work supplies a regulator with a measurable systemic quantity and three instruments. The intended impact is positive, namely making shared-model concentration legible as a stability risk. The foreseeable negative use is a firm reading the frontier as guidance on how much crowding it can add before a threshold binds, which we note is exactly the free-riding the public-good argument in Section 6 predicts.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No model, dataset or generative artifact with misuse potential is released. The work is closed-form theory plus a deterministic simulation harness.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: The order-flow simulator used for the panel 1 replication is prior published work and is cited; it is read-only throughout and unmodified. All other cited assets are published papers, cited in full in the references.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[N/A\]](#)

Justification: No new asset is released with this submission. The simulation harness and certificates are documented in the supplementary material and released on acceptance.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper involves no crowdsourcing and no research with human subjects.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper involves no research with human subjects, so no IRB review applies.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: No large language model is a component of the method. The contribution is closed-form analysis plus a numerical harness, and any LLM use was limited to writing and editing, which the guidelines exclude from this declaration.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.