
A Proof Attempt Told the Search Where to Look: A Case Study in Machine-Assisted Mathematical Research

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We report a research loop in which a numerical search and a human proof attempt
2 corrected each other, in enough detail that the mechanism is visible. An adversarial
3 machine search over experimental designs suggested that a certain estimation-
4 theoretic lower bound could not be beaten even by an adversary who knows the
5 model’s functional form. The best ratio it found after thousands of designs was
6 1.005, so we wrote the bound down as structure-proof. The proof we then wrote
7 did not cover that claim. It established the bound only within two curvature lengths
8 of the anchor point, and the obstruction was a single inequality that fails outside
9 that region. Searching exactly where the proof stopped produced a counterexample:
10 a four-point design with one far probe of weight 3×10^{-4} , attaining ratio 0.8517,
11 verified at 50-digit precision. The earlier verdict was a search artifact, and the
12 result is now two-sided and exact. Neither half of the loop found this alone. We
13 place the episode inside a larger body of machine-checked mathematics on the cost
14 of self-knowledge in performative prediction, in which eleven measurement-forced
15 pivots are recorded in the code where each occurred.

16 1 Introduction

17 This paper is about how a piece of mathematics got made, and it uses one concrete body of theory as
18 the specimen. Our answer to the venue’s guiding question is narrow and, we hope, checkable. In the
19 loop we ran, machines were good at refuting and at certifying, people were good at proving, and the
20 interesting events happened where the two disagreed. We document one such event completely, then
21 ask how far it generalizes using ten more of its own kind.

22 The specimen is a theory of what it costs a decision-maker to learn how its own deployed actions
23 reshape the world it then retrains on. Section 2 says what that means in plain terms. What matters
24 here is that the theory is large enough to have a texture: proved theorems, numerically supported
25 conjectures, and a live simulation pipeline, each of which turned out capable of correcting the others.
26 Over its development, measurement forced eleven changes to claims that had already been written
27 down, and each pivot is recorded in the code file where it happened.

28 We lead with the eleventh, because it is the one where the two halves of the loop needed each other.
29 A machine search believed a statement. A human proof attempt failed to reach the statement, and
30 failed in a way that named the exact region where it stopped. A second machine search, aimed only
31 at that region, refuted the statement.

32 **What this paper is not.** It is not an autonomous mathematical agent, and we make no claim in that
33 direction. The proofs here were written by people. What was automated is numerical checking against

declared tolerances, adversarial search over designs, and high-precision verification of candidate witnesses. Section 4 draws that line explicitly.

Contributions.

1. A reproducible account of one proof-and-search correction cycle, from the false belief through the failed proof to the frozen counterexample (Section 3). The resulting statement, Theorem 2, is two-sided and exact.
2. Ten further measurement-forced pivots from the same project, three of them the same pattern at smaller scale, and an account of what that sample does and does not support (Section 4).
3. The infrastructure that made the correction detectable: a theorem register carrying per-result proof status, preregistered numerical checks with tolerances declared before the run, and continuous integration over the derivation suite (Section 4), all shipped as an anonymized repository in the supplementary material.

2 The mathematical substrate, briefly

This section says what the theory is about, so that the counterexample in Section 3 has something to be a counterexample to. A reader who wants only the case study can take Theorem 1 on faith.

Consider a decision-maker whose deployed action changes the data it will later retrain on. A dealer who quotes a wide spread sees different order flow than one who quotes a tight spread, so the flow it must price is partly its own doing. This is performative prediction [Perdomo et al., 2020, Hardt and Mendler-Dünger, 2023]: retraining becomes a fixed-point iteration, and it converges to a point that is not the one maximizing value. Algorithms that close that gap [Izzo et al., 2021, Miller et al., 2021, Drusvyatskiy and Xiao, 2023] all need an estimate of how the distribution responds to the action, and they get it by perturbing deployments. Those perturbations are deployments of a real decision variable, so their cost lands in the decision-maker’s own objective.

Write h for the scalar action, h^* for the operating point the decision-maker would choose if it already knew everything, and $d_t = h_t - h^*$ for the deviation deployed at round t . Let ε be the response slope to be estimated, σ^2 the observation noise variance, Φ the value function, and γ_{PO} its curvature at h^* . The cost of exploring is the value forgone, $C_T = \sum_{t \leq T} \mathbb{E}[\Phi(h^*) - \Phi(h_t)]$. Formal statements are in Appendix A.

Theorem 1 (Exchange rate; T2). *Along any stationary exploration policy, for the least-squares slope estimator, $\text{Var}(\hat{\varepsilon}) \times C_T = \frac{1}{2} \gamma_{PO} \sigma^2 (1 + o(1))$, independent of the exploration amplitude, shape, schedule, and of how strongly the loop contracts.*

Learning how your own deployment moves the world costs a fixed amount. You can choose what to buy with it, but not how much to pay. The mechanism is a cancellation. Information accumulates as the centered energy of the design, cost accumulates as half the curvature times that same energy, and for mean-centered exploration the two coincide.

Three further results sit around that identity, all in Appendix A. No policy at all can beat the product, only match it (T3, T4). The cheapest way to spend a multidimensional budget is to probe where the objective is flat (T5a, C5.1). And a deployed agent can be made to correct only when a confidence sequence certifies that correcting cannot destabilize the retraining loop (L4, P7.1). The theory was also instantiated on a calibrated real-data problem and on a second economy, both in Table 1.

3 The lead story: a belief, a proof that missed it, and a witness

Theorem 1 is a statement about one estimator on one class of policies, and the minimax floors T3 and T4 extend it to all policies and all estimators. Both leave a natural question open. Suppose the adversary knows the functional form of the response, not just its slope. Concretely, suppose it knows that the response is $\tau_\theta(h) = C_0 + C_1 e^{-ch}$ and has only to learn the three parameters. Knowing the form is real information, so it might buy a design that beats the floor. Does it?

The machine’s answer, and why we believed it. We asked a search. For each candidate design μ the relevant quantity is the Cramér–Rao product normalized by the floor, $R(\mu) = (g^\top M(\mu)^{-1} g) \int (h - h^*)^2 d\mu$, where $M(\mu)$ is the design’s Fisher information and g is the gradient of the estimand. The




Figure 1: **Left:** the exchange rate of Theorem 1. Measured $\text{Var}(\hat{\varepsilon}) \times C_T$ against exploration amplitude at two contraction moduli, flat on the invariant (dashed). **Right:** the search that produced the false verdict. Best Cramér–Rao ratio with one tight probe at t , over the coarse weight grid with $w \geq 0.05$. On this family tight probes are monotonically worse, which is what we believed. The counterexample lives outside it: probe weight 3×10^{-4} , inside an asymmetric cluster the search never generated.

83 floor survives structural knowledge exactly when $R(\mu) \geq 1$ for every μ . Four thousand random
 84 designs plus 2,500 local refinements returned a minimum of $R = 1.0050$, and a systematic scan over
 85 designs pairing two near-anchor points with one distant probe found distant probes monotonically
 86 worse, reaching $R = 40.9$ at the grid’s furthest. Two independent strategies agreeing is the kind of
 87 evidence that gets written down, and we wrote it down: the floor is structure-proof, best attained ratio
 88 1.005.

89 **The proof did not reach the claim.** The proof attempt reduces $R(\mu)$ to a one-dimensional problem.
 90 A generalized Rayleigh identity gives $R(\mu) = \int d^2 d\mu / \inf \{ \int \phi^2 d\mu \}$, where $d(h) = h - h^*$ and
 91 the infimum runs over functions ϕ in the span of the model’s sensitivity components, normalized
 92 by $\phi'(h^*) = 1$. So it suffices to exhibit one such ϕ that is pointwise dominated by $|d|$. There is a
 93 canonical candidate, the unique element of the span with the right 2-jet at the anchor:

$$\phi_0(h) = \frac{2}{c} - e^{-c(h-h^*)} \left(\frac{2}{c} + h - h^* \right).$$

94 Writing $x = h - h^*$, the required domination $|\phi_0| \leq |x|$ holds for every $x \geq -2/c$, with equality
 95 only at $x = 0$ and $x = -2/c$. That is a clean elementary argument (Appendix B). It is also the whole
 96 story, because for $x < -2/c$ the domination does not merely resist proof. It is false, and it fails
 97 exponentially, since $\phi_0(x) = 2/c + e^{c|x|}(|x| - 2/c)$ there. The proof therefore delivered the floor on
 98 $[h^* - 2/c, \infty)$ and nowhere else, and it named the boundary.

99 **Searching where the proof stopped.** The failed step is not a gap in the argument. It is a location.
 100 We re-ran the adversarial search twice, once restricted to the reach and once with support allowed far
 101 below it. Inside the reach, a hardened search (three runs of 30,000 random designs plus 10,000 refine-
 102 ments) returned $R = 1.00054$, consistent with the proof and with an infimum of exactly 1. Allowed
 103 below the reach, the same search found $R = 0.85164$. The witness is a four-point design with support
 104 $\{1.8726, 1.8665, 1.3073, 0.0500\}$ and weights $\{0.003329, 0.955702, 0.040697, 0.000272\}$, in the
 105 register market with $c = 1.5$ and anchor $h^* = 1.8461$. Every support point except the one at 0.05 lies
 106 inside the reach, so that single below-reach probe of weight 3×10^{-4} is what breaks the floor. We
 107 froze the design and recomputed its ratio at 50-digit precision with `mpmath`, obtaining $0.851650\dots$;
 108 the check now runs on every push as V9.2.

109 **Why the first search missed it, exactly.** Two gaps, both now identified. The scan’s weight grid
 110 never went below 0.05, and the violating probe carries 3×10^{-4} . The random search did generate
 111 distant probes, but always alongside a symmetric near-anchor pair, and symmetric-pair-plus-far-probe
 112 families never break the floor: their ratio approaches 1 from above as the probe weight vanishes. The
 113 violation requires the asymmetric cluster. The earlier verdict was therefore neither noise nor a coding
 114 error. It was a correct summary of a search space that happened to exclude the answer, and nothing
 115 inside that search could have revealed the exclusion.

116 **Theorem 2** (Structure-proofness, exact reach; T9). (i) Every estimable design supported in $[h^* -$
 117 $2/c, \infty)$ satisfies $R(\mu) \geq 1$, strictly except on one two-point configuration, and $\inf_{\mu} R(\mu) = 1$.
 118 Within reach, knowing the family cannot beat the floor. (ii) The reach is exact: the four-point witness
 119 above attains $R = 0.8517$.

120 **The corrected result explains itself.** The two-sided statement says something the one-sided belief
 121 did not. Structure-proofness is an ignorance-of-curvature phenomenon with an exactly known extent.

122 A design reaching far enough below the anchor learns the curvature parameter c cheaply enough to
 123 bank that cheapness rather than pay for it, and the floor stops binding. Two boundary checks confirm
 124 the reading. With c known a priori, so that the ignorance is gone, the floor breaks outright and keeps
 125 falling with no limit the scan reaches ($R = 0.3002$ at probe $t = 1.0$, 0.0425 at $t = 0.05$, 0.0348 and
 126 still decreasing at $t = 0.001$). And the reach is exactly where a trust-region agent can deploy, since
 127 its clip is $r \leq 0.8 < 2/c = 1.333$: the constraint that keeps it safe is the one that keeps its lower
 128 bound valid. That was read off the corrected theorem, not designed.

129 4 The loop, and what it generalizes to

130 One episode is an anecdote. This section gives the sample it sits in and states what that sample
 131 supports.

132 **The division of labour.** Automated: numerical evaluation of every closed form against a simulation,
 133 at tolerances declared before the run; adversarial search over designs; high-precision arithmetic on
 134 frozen candidates; regression checking on every push. Not automated: every proof in this work, every
 135 choice of what to conjecture, and the decision to search below the reach. The machine in this loop is
 136 a fast, tireless, uncreative referee. It never proposed a theorem. It refuted several, and in the eleventh
 137 case only after a human argument told it where to look.

138 **Eleven pivots, three of them the same shape.** Measurement forced eleven changes to written-down
 139 claims, and three repeat the lead story at smaller scale. In E1, a saturation cap derived without noise
 140 failed by a factor of four against the noisy pipeline; the gap was a missing term, since retraining on
 141 noisy observations adds a stationary excitation floor of $(\sigma/\gamma)^2/(1 - m^2)$ per step, now itself verified
 142 within 0.2 to 6.4 percent. In E5, an injected misspecification ae^{-2ch} had decayed to about 10^{-3} by
 143 the operating point, so the anchored fit correctly never saw it; replacing it with a linear leak the family
 144 cannot represent produced the crossover the theory predicts. In E7, a claimed threefold advantage for
 145 shaped exploration collapsed to parity once the comparison arm was funded with the same energy,
 146 the unshaped arm having been given $r^2/3$ per step against the design's $2r^2/3$; the corrected claim
 147 is sharper, since the identification cliff is a support and amplitude phenomenon rather than a shape
 148 phenomenon. In each case the replacement was a stronger statement than the claim it displaced.

149 **What the sample supports.** A modest and specific thing: on this project, numerical falsification was
 150 cheap and frequent, and it improved claims rather than merely deleting them. It supports no claim
 151 about proof automation, about other areas of mathematics, or about what happens without a person
 152 deciding where to look. The eleventh pivot needed a proof to produce the region, and no further
 153 search inside the original space would have found it.

154 **The infrastructure, which is the transferable part.** A theorem register lists all 26 results with
 155 per-result proof status, distinguishing proved, proved for a stated class, and proved conditional
 156 on a hypothesis, so a numerically supported conjecture cannot quietly be cited as a theorem. A
 157 preregistered harness compares each measurement to a closed form at a tolerance fixed before the
 158 run; the state is 36 rows with 0 failures, one of them a deliberately out-of-scope cell reported as drift
 159 rather than excused. A derivation suite of 38 checks runs in continuous integration, and its first runs
 160 falsified three drafted claims, one with the wrong predicted sign. Every pivot is written into the code
 161 file where it occurred, which is why the eleventh survives as a record and not as a memory.

162 **The honesty discipline is part of the contribution.** We report the drift cell as drift. We label the
 163 sharp constant $\frac{1}{2}$ in T4 as numerically supported and open, since what is proved is the constant-factor
 164 version. And we say plainly that our own written verdict of structure-proofness was a search artifact.
 165 A record containing only successes cannot calibrate how much to trust the next numerical verdict,
 166 which is the only reason to keep one.

167 **Limitations.** Theorem 1 is local and its $o(1)$ is load-bearing: wide-side probing at finite amplitude
 168 undercuts the identity, and the out-of-scope saturating cell drifts +16 percent. The value-budget floor
 169 T4 is proved only to within a factor of 13.5 of the conjectured constant. Theorem 2(ii) lives inside
 170 the local-quadratic cost model, where under the true cost the same witness has ratio 1.0112, so the
 171 violation concerns the floor's own scope and not the underlying market. Which below-reach designs
 172 violate it, and by how much at most, is open. And this section rests on eleven pivots from one project
 173 by one group, which is a sample and not a study.

References

- Dmitriy Drusvyatskiy and Lin Xiao. Stochastic optimization with decision-dependent distributions. *Mathematics of Operations Research*, 48(2):954–998, 2023.
- Moritz Hardt and Celestine Mendler-Dünger. Performative prediction: Past and future. *arXiv preprint arXiv:2310.16608*, 2023.
- Zachary Izzo, Lexing Ying, and James Zou. How to learn when data reacts to your model: Performative gradient descent. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.
- John Miller, Juan C. Perdomo, and Tijana Zrnic. Outside the echo chamber: Optimizing the performative risk. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.
- Vignesh Nagarajan and Shriraghav Ashok. REFLEX: Reflexive equilibrium fixed-point learning for endogenous exchanges. *arXiv preprint arXiv:2608.16155*, 2026.
- Juan C. Perdomo, Tijana Zrnic, Celestine Mendler-Dünger, and Moritz Hardt. Performative prediction. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020.

A Assumptions and statements

This appendix collects the standing assumptions and the formal statements of every result the body refers to, and nothing else. Identifiers (L1, T2, T9, and so on) are the register identifiers used throughout the supplementary repository, so each statement here can be matched to its derivation document, its proof, and its numerical checks. The register itself, with per-result proof status, is `mlxor-derivations/THEOREMS.md`.

A.1 Standing assumptions

A1 (local quadratic scope): Φ is C^3 and strongly concave on the operating interval, with expansions taken at h^* and third-order remainder bounded by $L_3 = \sup |\Phi'''|$. **A2** (exploration class): all policies are non-anticipating; A2-sym adds $\mathbb{E}[u_t | \mathcal{F}_t] = 0$, A2-adaptive is unrestricted, and costs are measured from the certainty-equivalent anchor. **A3** (noise exogeneity): the observation noise ζ_t is i.i.d., mean zero, variance σ^2 , independent of $\mathcal{F}_t$. A3' relaxes this to a feedback channel with gain ϕ . **A4** (trust region): $|d_t| \leq r$ pathwise. **A5** (stationary scale): $S_T = \Theta(T)$. **A6** (stable base loop): the retraining modulus satisfies $m < 1$.

The local dynamics are $d_{t+1} = -m d_t + u_t$ with $d_t = h_t - h^*$, observations $\tau_t = \tau(h^*) - \varepsilon d_t + \zeta_t$, and Φ satisfying $\Phi'(h^*) = \delta_0$ and $-\Phi'' = \gamma_{\text{PO}}$ at the operating point. S_{xx} denotes the centered design energy $\sum_{t \leq T} (d_t - \bar{d})^2$.

A.2 The exchange rate and its floors

Lemma 1 (Cost equivalence; L1). *Under A1–A3 with conditionally symmetric exploration, the expected incremental performative cost satisfies $C_T = \frac{1}{2} \gamma_{\text{PO}} \mathbb{E}[\sum_{t \leq T} d_t^2] + R_T$ with $|R_T| \leq \frac{L_3}{6} \mathbb{E} \sum_t |d_t|^3$, and the realized cost's variance decomposes exactly as $\text{Var}(C_T^{\text{real}}) = \delta_0^2 \text{Var}(\sum_t d_t) + \frac{\gamma_{\text{PO}}^2}{4} \text{Var}(\sum_t d_t^2)$ with vanishing cross term.*

Theorem 3 (Information saturation; T1). *With no deliberate exploration and $\rho(M) < 1$, the total design energy of the retraining trajectory equals $d_0^\top P d_0$, where P solves the discrete Lyapunov equation $P = I + M^\top P M$. Consequently the Fisher information for the response Jacobian is bounded in every direction, uniformly in the horizon. In the scalar case the cap is $d_0^2/(1 - m^2)$, strictly increasing in m : a fast-contracting loop reveals least about its own performativity (C1.1). Under the noisy feedback channel A3' the trajectory additionally carries a stationary excitation floor of $(\sigma/\gamma)^2/(1 - m^2)$ per step.*

The body's Theorem 1 is the asymptotic form of the following pathwise identity, which is what is proved in Appendix B.

Theorem 4 (The exchange rate, pathwise; T2). *Under A1, A2-sym, A3 and A6, in the stationary regime, on every realization,*

$$\text{Var}(\hat{\varepsilon} \mid \text{design}) \times \frac{1}{2} \gamma_{\text{PO}} \sum_{t \leq T} d_t^2 = \frac{1}{2} \gamma_{\text{PO}} \sigma^2 \left(1 + \frac{T \bar{d}^2}{S_{xx}} \right) = \frac{1}{2} \gamma_{\text{PO}} \sigma^2 (1 + O_P(1/T)),$$

independent of the exploration intensity, the horizon, and the modulus.

Theorem 5 (Minimax lower bound, deviation budget; T3). *Over all non-anticipating exploration policies with $\sum_{t \leq T} d_t^2 \leq D$ almost surely and all estimators, biased included, the minimax risk over an interval of width w satisfies $\inf \sup_{\varepsilon} \mathbb{E}[(\hat{\varepsilon} - \varepsilon)^2] \geq \sigma^2/(D + \sigma^2(\pi/w)^2)$, so the exchange-rate product is a floor, attained with equality by symmetric probing with least squares.*

Lemma 2 (Exploitation and information; L2). *For the symmetric two-point response family with half-width δ and any non-anticipating policy under the certainty-equivalent anchor, $\mathbb{E}[\text{gain}_T] \leq g_1 \delta^{3/2} T^{1/2} (\mathbb{E} S_T)^{3/4} \sigma^{-1/2}$ with $g_1 = \beta |h^* - \psi|$, via the identity $\mathbb{E}[\mathbb{E}[s | \mathcal{F}_t]] = \text{TV}(P_t^+, P_t^-)$ and Pinsker's inequality. At the minimax scale this is an $O(T^{-1/2})$ fraction of the exploration cost.*

Theorem 6 (Minimax lower bound, value budget; T4). *Under stationary-scale exploration and the certainty-equivalent anchor, for every non-anticipating policy and every estimator, $\sup \text{Var}(\hat{\varepsilon}) \times C_T \geq \gamma_{\text{PO}} \sigma^2/27 (1 - O(T^{-1/2}))$, the supremum over the two-point response family at the minimax scale, the constant arising from a Le Cam reduction optimized at $c = 2/3$. The sharp constant $\frac{1}{2}$,*

235 against which no simulated policy fell by more than Monte Carlo error, is conjectured and is part of
 236 the open problem OPEN-1. It is not proved here or anywhere in the repository.

237 A.3 Design and certification

238 **Theorem 7** (A-optimal exploration; T5a). *Under the value budget $\text{tr}(\Gamma_{\text{PO}} M) \leq B$ on the exploration*
 239 *second moment M , the A-optimal design is $M^* = (B / \text{tr} \Gamma_{\text{PO}}^{1/2}) \Gamma_{\text{PO}}^{-1/2}$, with value $(\text{tr} \Gamma_{\text{PO}}^{1/2})^2 / B$:*
 240 *explore where the objective is flat.*

241 **Corollary 1** (Price of isotropic jitter; C5.1). *Isotropic exploration at matched budget overpays the*
 242 *A-optimal design by exactly the curvature-dispersion factor $F = d \text{tr}(\Gamma_{\text{PO}}) / (\text{tr} \Gamma_{\text{PO}}^{1/2})^2 \geq 1$, with*
 243 *equality if and only if all curvature directions are equal.*

244 **Theorem 8** (Chebyshev unidentifiability; T6). *The sensitivity system of the structural family $\tau(h) =$*
 245 *$C_0 + C_1 e^{-ch}$ is an extended Chebyshev system. Any design with fewer than three support points has*
 246 *singular Fisher information, and since noiseless retraining trajectories have at most two effective*
 247 *support points, no retraining trajectory at any modulus identifies the response curvature c .*

248 **Theorem 9** (Anchoring crossover; T7). *Under the value budget B and horizon T , the*
 249 *budget-and-horizon-optimal nonparametric secant estimator has $\text{MSE}_{\text{np}} = \gamma_{\text{PO}} \sigma^2 / (2B) +$*
 250 *$(\tau'''(h^*)/6)^2 (2B / (\gamma_{\text{PO}} T))^2$, so anchoring to a structural family with misspecification δ_{mis} wins, to*
 251 *leading order, if and only if $\delta_{\text{mis}} < |\tau'''(h^*)| B / (3 \gamma_{\text{PO}} T)$.*

252 **Lemma 3** (Perturbed modulus; L4). *If the corrected update uses an estimated response with C^1*
 253 *error e , the closed-loop modulus satisfies $|\hat{\rho} - \rho| \leq \eta \beta (\|e\|_\infty + \|h^* - \psi\| \|e'\|_\infty)$.*

254 **Proposition 1** (Safety under pessimism; P7.1). *Given an anytime-valid confidence sequence for the*
 255 *response parameters at level $1 - \delta$, the pessimism rule, which keeps the worst-case modulus over the*
 256 *confidence set below $1 - c_{\text{margin}}$ and freezes the correction otherwise, keeps the realized modulus*
 257 *within the margin for all t simultaneously with probability at least $1 - \delta$. The result is conditional on*
 258 *the validity of the confidence sequence, and the register records it as such.*

259 A.4 Structure-proofness

260 **Theorem 10** (Structure-proofness of the floor, exact reach; T9). *Fix the structural family $\tau_\theta(h) =$*
 261 *$C_0 + C_1 e^{-ch}$ ($C_1 > 0$, $c > 0$), the anchor $h^* > 0$, the sensitivity $s(h) = \partial_\theta \tau_\theta(h)$, and the*
 262 *estimand gradient $g = \partial_\theta \varepsilon(h^*) = -s'(h^*)$. For a finitely supported design μ define the normalized*
 263 *Cramér–Rao product*

$$R(\mu) = (g^\top M(\mu)^{-1} g) \int (h - h^*)^2 d\mu, \quad M(\mu) = \int s s^\top d\mu,$$

264 *so that $R(\mu) \geq 1$ states that the family-knowing Cramér–Rao product $\text{Var}(\hat{\varepsilon}) \times C_T$ is at least*
 265 *the exchange-rate floor $\frac{1}{2} \gamma_{\text{PO}} \sigma^2$ under the local-quadratic (AI) cost, for unbiased estimators and,*
 266 *in the local-asymptotic-minimax sense, for all regular estimators. (i) Within reach the floor is*
 267 *structure-proof: if $\text{supp}(\mu) \subset [h^* - 2/c, \infty)$, then either $\varepsilon(h^*)$ is not estimable under μ , in which*
 268 *case the bound is infinite, or $R(\mu) \geq 1$, strictly whenever μ places mass off the two equality points*
 269 *$\{h^*, h^* - 2/c\}$; and $\inf_\mu R(\mu) = 1$, approached but not attained by symmetric three-point designs*
 270 *collapsing at h^* . In particular the floor binds on every trust-region design with $r \leq 2/c$. (ii)*
 271 *The reach is exact: the pointwise domination underlying (i) fails strictly below $h^* - 2/c$, and an*
 272 *explicit four-point design with a single below-reach probe of weight 3×10^{-4} attains $R = 0.8517$,*
 273 *computer-verified at 50-digit precision as check V9.2.*

274 B Proofs

275 This appendix proves the two results the body argues from: the exchange rate of Theorem 1, in
 276 the exact pathwise form (Theorem 4), and the two-sided structure-proofness result of Theorem 2
 277 (Theorem 10). Both proofs are reproduced from the derivation repository without change; the check
 278 identifiers V2.1 and V9.1 to V9.4 refer to its verification suite. Proofs of the remaining statements in
 279 Appendix A are in that repository and are not reproduced here.

280 **The cost-equivalence lemma**

281 *Proof of Lemma L1 (cost equivalence).* Taylor's theorem with Lagrange remainder gives, for each t ,

$$\Phi(h^*) - \Phi(h_t) = -\delta_0 d_t + \frac{1}{2}\gamma_{\text{PO}} d_t^2 + r_t, \quad |r_t| \leq \frac{L_3}{6} |d_t|^3,$$

282 with $L_3 = \sup |\Phi'''|$ on the operating interval. Summing over $t \leq T$ and taking expectations, it
 283 remains to show $\mathbb{E}[\sum_t d_t] = 0$ under A2-sym. The recursion gives $\mathbb{E}[d_{t+1} | \mathcal{F}_t] = -m d_t + \mathbb{E}[u_t |$
 284 $\mathcal{F}_t] = -m d_t$, so $\mathbb{E}[d_{t+1}] = -m \mathbb{E}[d_t]$ by the tower property, and $\mathbb{E}[d_t] = (-m)^t \mathbb{E}[d_0] = 0$ when
 285 the loop starts at the operating point or in the stationary regime. The quadratic term contributes
 286 $\frac{1}{2}\gamma_{\text{PO}} \mathbb{E}[\sum_t d_t^2]$, which is the claim.

287 For the variance decomposition, write the realized cost (net of the cubic remainder) as $C =$
 288 $-\delta_0 X + \frac{1}{2}\gamma_{\text{PO}} Y$ with $X = \sum_t d_t$, $Y = \sum_t d_t^2$. Then $\text{Var}(C) = \delta_0^2 \text{Var}(X) + \frac{\gamma_{\text{PO}}^2}{4} \text{Var}(Y) -$
 289 $\delta_0 \gamma_{\text{PO}} \text{Cov}(X, Y)$. Under Gaussian innovations the vector $(d_1, \dots, d_T)$ is jointly Gaussian with
 290 mean zero, and every third moment of a zero-mean Gaussian vector vanishes; since $\text{Cov}(X, Y) =$
 291 $\sum_{t,s} \mathbb{E}[d_t d_s^2]$ is a sum of third moments, it is zero. (For non-Gaussian symmetric innovations the
 292 same conclusion follows from the sign symmetry $d \mapsto -d$ of the stationary law.) Verified: V0.1–V0.3,
 293 including the necessity check that an asymmetric rule breaks the expectation identity by exactly
 294 $-\delta_0 \mathbb{E}[\sum_t d_t]$. $\square$

295 **The exchange rate**

296 *Proof of Theorem T2 (pathwise exchange rate).* Under A3, conditionally on the design $\{d_t\}$ the ob-
 297 servations are independent with means linear in ε and variance σ^2 , so the least-squares slope has
 298 conditional variance $\text{Var}(\hat{\varepsilon} | \text{design}) = \sigma^2 / S_{xx}$ (Gauss–Markov). Multiplying by the quadratic
 299 cost $\frac{1}{2}\gamma_{\text{PO}} \sum_t d_t^2$ and using $\sum_t d_t^2 = S_{xx} + T\bar{d}^2$ gives the displayed identity exactly, on ev-
 300 ery realization. In the stationary regime $\bar{d} = O_P(T^{-1/2})$ by the stationary long-run variance
 301 $\text{Var}(\sum_{t \leq T} d_t) = Tv(1-m)/(1+m)(1+o(1))$, while $S_{xx}/T \rightarrow v$ a.s., so $T\bar{d}^2/S_{xx} = O_P(1/T)$.
 302 Verified: V2.1 (machine precision). $\square$

303 **Structure-proofness, and the exactness of the reach**

304 *Proof of Theorem T9 (structure-proofness within reach).* Write $x = h - h^*$, $d(h) = x$, and $V =$
 305 $\text{span}\{1, e^{-ch}, he^{-ch}\}$, the span of the components of s (a bijective correspondence $u \leftrightarrow \phi_u = u^\top s$
 306 between $\mathbb{R}^3$ and V , since $C_1 \neq 0$). Differentiating s gives $s'(h) = (0, -ce^{-ch}, C_1 e^{-ch}(ch - 1))^\top$,
 307 so $g = -s'(h^*)$ componentwise, which is the stated estimand gradient.

308 *Step 1: one-dimensional reduction.* For $M(\mu) \succ 0$ the generalized Rayleigh identity gives

$$g^\top M(\mu)^{-1} g = \sup_{u \neq 0} \frac{(u^\top g)^2}{u^\top M(\mu) u} = \sup_{\phi \in V} \frac{\phi'(h^*)^2}{\int \phi^2 d\mu} = \left(\inf \left\{ \int \phi^2 d\mu : \phi \in V, \phi'(h^*) = 1 \right\} \right)^{-1},$$

309 using $u^\top M u = \int \phi_u^2 d\mu$ and $u^\top g = -\phi'_u(h^*)$ (the sign is lost in the square). Hence

$$R(\mu) = \frac{\int d^2 d\mu}{\inf \left\{ \int \phi^2 d\mu : \phi \in V, \phi'(h^*) = 1 \right\}}.$$

310 *Step 2: the candidate and its pointwise domination.* Let

$$\phi_0(h) = \frac{2}{c} - e^{-cx} \left(\frac{2}{c} + x \right) = \frac{2}{c} \cdot 1 - e^{ch^*} \left(\frac{2}{c} - h^* \right) e^{-ch} - e^{ch^*} (he^{-ch}) \in V,$$

311 the unique element of V with 2-jet $(0, 1, 0)$ at h^* (the 2-jet map is invertible on V : in the basis
 312 $\{1, e^{-ch}, he^{-ch}\}$ its determinant is $c^2 e^{-2ch^*} \neq 0$); the expansion is $\phi_0(x) = x - (c^2/6)x^3 + O(x^4)$.
 313 Writing $f(x) = \frac{2}{c} - e^{-cx}(\frac{2}{c} + x)$ we claim

$$|f(x)| \leq |x| \quad \text{for all } x \geq -2/c, \quad \text{with equality iff } x \in \{0, -2/c\},$$

314 and $|f(x)| > |x|$ strictly for every $x < -2/c$. Indeed: (a) $v(x) = f(x) - x$ has $v(0) = 0$ and
 315 $v'(x) = e^{-cx}(1 + cx) - 1 < 0$ for $x \neq 0$ (since $1 + t < e^t$ for $t = cx \neq 0$); $v' \leq 0$ with equality
 316 only at the isolated point 0 makes v strictly decreasing on $\mathbb{R}$, so $f > x$ on $x < 0$ and $f < x$ on

317 $x > 0$. (b) $f(0) = 0$ and $f'(x) = e^{-cx}(1 + cx) > 0$ for $x \geq 0$, so $f \geq 0$ on $[0, \infty)$ and, with (a),
318 $0 \leq f \leq x$ there. (c) $q(x) = f(x) + x$ has $q(-2/c) = 0 = q(0)$ and $q'(x) = e^{-cx}(1 + cx) + 1$,
319 where $r(x) = e^{-cx}(1 + cx)$ is strictly increasing on $x < 0$ ($r'(x) = -c^2 x e^{-cx} > 0$) from
320 $r(-2/c) = -e^2 < -1$ to $r(0) = 1$; so q' changes sign exactly once on $(-2/c, 0)$, from negative
321 to positive, and $q \leq 0$ there with equality only at the endpoints; with (a) ($f \geq x$, equality only
322 at 0) this gives $x \leq f \leq -x$ on $[-2/c, 0]$. (d) For $x < -2/c$, $r(x) < -e^2$ gives $q' < 0$, so
323 moving left from $-2/c$ the function q increases above 0: $f(x) > -x = |x| > 0$, and in fact
324 $f(x) = 2/c + e^{c|x|}(|x| - 2/c)$ grows exponentially.

325 *Step 3: the bound.* If $\text{supp}(\mu) \subset [h^* - 2/c, \infty)$, Step 2 gives $\int \phi_0^2 d\mu \leq \int d^2 d\mu$, so by Step 1,
326 $R(\mu) \geq \int d^2 d\mu / \int \phi_0^2 d\mu \geq 1$. Strictness: equality in Step 2 holds only at $x \in \{0, -2/c\}$, so any
327 μ placing positive mass off $\{h^*, h^* - 2/c\}$ (in particular any with three distinct support points)
328 has $\int \phi_0^2 d\mu < \int d^2 d\mu$ and $R(\mu) > 1$. If $M(\mu)$ is singular, two cases: $g \notin \text{range } M(\mu)$ makes
329 $\varepsilon(h^*)$ non-estimable (infinite bound); $g \in \text{range } M(\mu)$, which does occur for special two-point
330 designs (a one-parameter family of in-reach pairs; their finite pseudo-inverse products were checked
331 to obey the bound, e.g. $R = 1.063$ at the pair $\{1.627, 2.127\}$), keeps the sup-representation valid
332 over $\{u : u^\top M u > 0\}$, and $\int \phi_0^2 d\mu = 0$ would put $u_0 \in \text{null } M$ with $u_0^\top g = -1 \neq 0$, contradicting
333 $g \in \text{range } M$; so $\int \phi_0^2 d\mu > 0$ and the bound goes through verbatim. Throughout, the Cramér–Rao
334 functional bounds unbiased estimation, and all regular estimators in the local-asymptotic-minimax
335 sense.

336 *Step 4: the infimum is 1.* The Cramér–Rao functional is invariant under C^1 reparametriza-
337 tion, and the 2-jet map $\theta \mapsto (\tau_\theta(h^*), \varepsilon(h^*), \tau_\theta''(h^*))$ is a local diffeomorphism: its Jacobian
338 $[s(h^*) - s'(h^*) \ s''(h^*)]$ has determinant $C_1 c^2 e^{-2ch^*} \neq 0$ (direct computation; V9.4). In
339 jet coordinates η the model reads $\tau = \eta_1 - \eta_2 x + \frac{1}{2} \eta_3 x^2 + O(x^3)$ with sensitivity $\tilde{s}(x) =$
340 $(1, -x, x^2/2)^\top + O(x^3)$, and the estimand is η_2 . For the symmetric three-point design μ_δ uniform
341 on $\{h^* - \delta, h^*, h^* + \delta\}$, scale by $D = \text{diag}(1, \delta, \delta^2)$: $M_\eta(\mu_\delta) = D \hat{H}_\delta D$ with $\hat{H}_\delta \rightarrow \hat{H}_0$ invertible
342 (the exact quadratic-model moment matrix), so $(M_\eta^{-1})_{22} = \delta^{-2} (\hat{H}_\delta^{-1})_{22}$ and $R(\mu_\delta) \rightarrow R_{\text{quad}}$. In
343 the exact quadratic model, for *any* symmetric design (odd moments $m_1 = m_3 = 0$), the $(2, 2)$
344 cofactor gives $(M_\eta^{-1})_{22} = 1/m_2$ and $R_{\text{quad}} = m_2 \cdot (1/m_2) = 1$. Hence $R(\mu_\delta) \rightarrow 1$ (measured rate
345 $R(\mu_\delta) - 1 = O(\delta^2)$; V9.2), the infimum is 1, and by Step 3 it is not attained.

346 *Step 5: the reach is exact.* Below $h^* - 2/c$ the domination of Step 2 fails strictly, and the fail-
347 ure is realized by a design, not only by the candidate: the frozen four-point witness (support
348 $\{1.8726, 1.8665, 1.3073, 0.0500\}$, weights $\{0.003329, 0.955702, 0.040697, 0.000272\}$ after normal-
349 ization, in the register market $c = 1.5$, frozen anchor $h^* = 1.8461$) has $R = 0.851650 \dots$, verified
350 at 50-digit precision (mpmath) and reproduced in float64 by `run_open1.py` (V9.2). All its support
351 except the probe at 0.05 lies within reach: one vanishing-weight below-reach probe breaks the floor.
352 Symmetric pair-plus-far-probe families do *not* break it (their ratio approaches 1 from above as the
353 probe weight vanishes); the violation requires the asymmetric cluster, which is why coarse searches
354 missed it. Verified: V9.1–V9.4. $\square$

355 C Verification rows

356 D Reproducibility and scope of this version

357 **What ships.** The supplementary repository contains the derivation documents and the theorem register
358 (26 results, 38 numerical derivation checks), the simulation pipeline (environments, estimators,
359 designs, agents, baselines; 36 measured-versus-predicted rows; 9 unit tests), the real-data port with its
360 10-cell validation, and the eleven-pivot record in the pipeline README. Both verification suites run
361 in continuous integration on every push in the project repository; the two workflow files are not part
362 of this anonymized copy. The structural market used throughout, and the corporate-bond calibration
363 behind the real-data row of Table 1, are those of REFLEX [Nagarajan and Ashok, 2026], on which
364 this work builds. All results are deterministic from seeds, numpy-only, and CPU-only. Reported final
365 errors are 300-step tail averages over the deployed path, stable within 0.01 under horizon perturbation,
366 with a worst measured spread of 0.0062.

367 **What is stated here but not proved here.** This version proves Theorem 1 and Theorem 2, the two
368 results the body argues from. The minimax floors and the design and certification results (T3, T4,

Table 1: Selected rows from the preregistered check harness. The full table (36 rows, 0 FAIL), all seeds, tolerances, and code are in the supplementary repository. The drift row is a deliberately out-of-scope cell, measured and reported rather than excused.

Claim (register)	Measured	Predicted	
Exchange rate, $m=0.3$, two amplitudes (T2)	0.01062, 0.01070	0.01056	pass
Exchange rate, $m=0.6$, two amplitudes (T2)	0.00549, 0.00569	0.00535	pass
Saturating environment (outside A1)	+16% over rate	out of scope	drift
Feedback floor, $m=0.3/0.6/0.85$ (T1 corollary)	within 0.2–6.4%	$(\sigma/\gamma)^2/(1-m^2)$	pass
Design arm $\text{sd}(\hat{c})$; jitter/design ratio (T6)	0.234; $7.8\times$	Fisher 0.248; ≥ 3	pass
Certified agent reaches h_{PO} ; trust region (L4)	1.258; 0 violations	1.296; 0	pass
Isotropic/A-optimal risk ratio, $d=8$ (T5a)	1.529	$F = 1.506$	pass
Nonparametric MSE at (B, T) optimum (T7)	0.0132	0.0123	pass
Second economy, 4 cells (T2, exact in LQ); agent	0.384, 0.192; 12.65	0.384, 0.192; 12.5	pass
Multi-book shaped/isotropic steady-state risk (T5a)	1.20 (flat control 1.00)	$> 1.15 (\approx 1)$	pass
Real-data calibration port, 10 cells	10/10 reproduce	within 0.8%	pass
Within-reach min CR ratio; witness beyond (T9)	1.0005; 0.8517	≥ 1 ; < 1	pass

369 T5a, C5.1, T6, T7, L4, P7.1) are stated in Appendix A and proved in the supplementary repository, in
370 `mlxor-derivations/latex/proofs.tex` and the corresponding derivation documents, with the
371 register recording each result’s proof status.

372 **Provenance of the numbers.** Every quantity in the body and in Table 1 is copied from a committed
373 artifact: `results/RESULTS.md` for the harness rows, `results/OPEN1.md` for the reach, the witness,
374 and the search history, `results/REALDATA.md` for the calibration port, `results/STABILITY.md`
375 for the tail-average stability, `results/full_run.log` for quantities not surfaced in the summary
376 tables, and `THEOREMS.md` for the register. The runs behind them were executed and recorded on 1
377 September 2026. This version re-narrates them and does not re-run them, since none of its changes
378 touch their inputs.